

Evidence maturity — what is executed, what is same-context, what is still pending

● EXECUTED

mechanics + known-answer contracts

- epistemic-noninterference tests + integration-path mutations (executable)
- hostile deterministic known-answer worlds behave as specified
- RAKLV3State carries immutable ScientificAuthorityProjection; transitions typed by ownership
- v1 fingerprint byte-stable; explicit v2 fingerprint includes scientific authority
- SEARCH·JUMP·GLUE·LIFT router runs (proposal-only)

● SAME-CONTEXT-ONLY

self-repair / development-time

- v1.2 packet: 12/12 schema-valid native records, intended $S_0 \rightarrow S_n$ chronology
- uncontaminated fresh-transfer construction
- interface self-repair after two *preserved* prompt-interface failures
- same-session critique is useful signal

useful, but not independent review and not a superiority result

● PENDING / BLOCKED

empirical headline reads **PENDING**

- matched-workflow superiority (LEARNING vs RESET) — 0 registered successes @0.5B–7B
- fresh-assurance self-evolution gain Δ_A — CANNOT_IDENTIFY
- real quant-finance / domain result — preregistered
- fair external-LLM comparator
- independent external review (#216)

Master blocker. MODEL_CAPABILITY_FLOOR_7B(_V2_EXEC) with CAPABLE_MODEL_AVAILABLE = NO_REFUTED: the empirical superiority headline is **pending**, not refuted. Architecture/integration invariants are executed; efficacy, safety and cost require a capable model and separately registered experiments.